

Group: LSC Group 25

Members: Kozy-Korpesh Tolep (S5302354) and Urwa Fatima (S5156692)

Motivation:

The amount of data has risen exponentially in the last couple of years, distributed system technologies like Spark are overcoming the challenges of handling and storing data with efficient and optimized APIs. Therefore this lab project is developed to learn more about Spark SQL Dataframe API that handles advanced and complex scenarios for analyzing data.

Objective:

The objective is to implement complex transformations, selections, and extract information from a large [dataset](#) of BigMac Burger prices over several years [1]. This dataset has 812 rows and 14 columns.

Following questions are addressed:

1. Find top three countries with most expensive BigMac in 2011
2. Find top 20 countries with max price for BigMac over the years
3. Find the difference between maximum value and minimum value for dollar exchange for each country and sort by difference in descending order
4. Show average local price of Bigmac for each year and find the difference with the previous year's average local price and sort difference in descending order
5. For each country show largest increase in dollar exchange rate(dollar_ex) comparing month of each year and sort in descending order
6. For each year give a rank for countries average dollar_price for BigMac.

Saved successfully!



API:

A DataFrame is a distributed collection of data with well defined rows and columns which makes it conceptually equivalent to a table in a relational database or a dataframe in R or Python. It differs in executional efficiency because of the SQL-like optimizer called Catalyst working behind the scenes. DataFrame is immutable and it is structured to have lazily evaluated plans that specify what operations to apply to data in order to generate some output.

Sources of data in DataFrames: DataFrames can be constructed from a wide array of sources such as structured data files, tables in Hive, external databases, or existing RDDs. The DataFrame API is available in Scala, Java, Python, and R.

The DataFrame API can act as a distributed SQL query engine as the input data can be queried by using ad-hoc methods and an SQL-like language for structured data manipulation. Spark has mainly two types of method which can be performed on DataFrame. These methods also referred as operations are either Action or Transformation

Transformation: The Spark operations that reads a DataFrame, manipulates some of the columns, and returns another DataFrame are considered transformational operations. Functions like `select()`, `filter()` and `groupBy()` are example of transformation

Action: The Spark operations that either returns a result or writes to the disc are considered as action function for example, `count()` and `collect()`.

For a complete list of the types of operations that can be performed on a DataFrame can be referred through the [API Documentation](#).

As DataFrames has the feature of lazily evaluation. This means when Spark transforms data, it does not immediately compute the transformation but plans on how to compute later. When actions such as `collect()` are explicitly called, then computation starts.

▼ Difference between RDD vs Dataframes vs Datasets:

This table shows the similarities and differences of these Spark data manipulation APIs [3].

RDD	Dataframes	Datasets
Fault Tolerant	Fault Tolerant	Fault Tolerant
Distributed	Distributed	Distributed
Immutable	Immutable	Immutable
No Schema	Schema	Schema
Slow on Non-JVM languages	Faster	Faster
No Execution Optimzation	Catalyst Optimization	Optimization
Low Level	High Level	High Level
No SQL Support	SQL Support	SQL Support
Type Safe	No Type Safe	Type Safe
Syntax Error Detected at Compile time	Syntax Error Detected at Compile time	Syntax Error Detected at Compile time
Analysis Error at Compile Time	Analysis Error Detected at Run Time	Analysis Error at Compile Time
JAVA, SCALA, Python, R	JAVA, SCALA, Python, R	JAVA, SCALA

```
#run this block if the pyspark library is not installed
!pip install pyspark
```

Collecting pyspark

Downloading pyspark-3.2.0.tar.gz (281.3 MB)


```

root
|-- date: string (nullable = true)
|-- iso_a3: string (nullable = true)
|-- currency_code: string (nullable = true)
|-- name: string (nullable = true)
|-- local_price: double (nullable = true)
|-- dollar_ex: double (nullable = true)
|-- dollar_price: double (nullable = true)
|-- GDP_dollar: double (nullable = true)
|-- adj_price: double (nullable = true)
|-- USD: double (nullable = true)
|-- EUR: double (nullable = true)
|-- GBP: double (nullable = true)
|-- JPY: double (nullable = true)
|-- CNY: double (nullable = true)

```

```
print(datasetDF.count(), len(datasetDF.columns))
```

```
812 14
```

Question 1:

Find top three countries with most expensive BigMac in 2011.

Solution:

First we need to filter and leave only records which were made in 2011. However, the date column as it's in the format - 'yyyy-mm-dd', we need to use date() function to obtain a year from there. After that, we select only the name of the country and sort descending by 'dollar_price' and show top 3.

```

# select() returns a new dataframe with the provided column name and where() filters the data
# sort() the function arranges data in the column specified in ascending order by default.

```

```
data = datasetDF.where(year('date') == 2011).select('name').sort(desc('dollar_price'))
```

Saved successfully!



```

+-----+
|      name|
+-----+
|      Norway|
|Switzerland|
|      Sweden|
+-----+

```

only showing top 3 rows

CPU times: user 16.2 ms, sys: 2.42 ms, total: 18.7 ms

Wall time: 1.05 s

Question 2:

Find average dollar_price for BigMac for each year and for all countries and sort by year.

Solution:

To find average 'dollar_price' over the years for each country we use the aggregation function avg() from pyspark.sql and group records by 'year'. After that we sort data by 'year' column.

```
# selectExpr() is a variation of select() that projects a set of SQL expression
# .agg() function compute aggregates and return dataframe as a result. The available aggregat

data = datasetDF.selectExpr('year(date) as year', 'dollar_price').groupBy('year').agg(F.avg('
data.show()
```

```
+-----+-----+
|year|      avg_price|
+-----+-----+
|2011|3.9364215199569053|
|2012| 3.64763559762034|
|2013|3.8091063435738315|
|2014| 3.836926679387834|
|2015|3.5356402978675883|
|2016| 3.373439669452451|
|2017| 3.541480205052853|
|2018|3.7087346670848276|
|2019| 3.60018134836105|
|2020| 3.675629022124684|
|2021|3.7285290062520344|
+-----+-----+
```

```
CPU times: user 28.4 ms, sys: 1.75 ms, total: 30.1 ms
Wall time: 2.27 s
```

Question 3:

Saved successfully!



value and minimum value for dollar exchange for each
ending order

Solution:

For solving this question we need to get two aggregation functions from one column. That is why we imported pyspark.sql to use max() and min() aggregation functions to one column. After that, we need alias function to name the new columns created by aggregation fuctions. We need to change it from 'max(dollar_ex)' and 'min(dollar_ex)' because when we will use raw sql 'max(dollar_ex)' column can be interpreted by selectExpr() function as a new max() aggregation function. After naming two columns obtained from aggregation functions we use selectExpr() function and subtract minimum of dollar exchange from maximum and sort by their difference.

```
data = datasetDF.groupBy('name').agg(F.max('dollar_ex').alias('max_dollar_ex'), F.min('dollar_ex').alias('min_dollar_ex')).selectExpr('name', 'max_dollar_ex - min_dollar_ex as difference_of_dollar_ex').show()
```

name	difference_of_dollar_ex
Indonesia	5994.5
Colombia	2071.61
Chile	324.95000000000005
South Korea	177.20000000000005
Hungary	123.30555000000001
Argentina	92.20124999999999
Pakistan	80.155
Japan	47.015
Russia	46.9025
Vietnam	35.5
India	30.797500000000007
Philippines	12.887
Egypt	12.808
Mexico	10.76525
Sri Lanka	10.5
South Africa	9.903649999999999
Czech Republic	8.789850000000001
Turkey	6.8344
Thailand	6.469999999999999
Taiwan	5.2514999999999965

only showing top 20 rows

CPU times: user 21.3 ms, sys: 0 ns, total: 21.3 ms

Wall time: 1.24 s

Question 4:

Show average local price of Bigmac for each year and find the difference with the previous year's average local price and sort difference in descending order

Saved successfully!



Firstly, we need to find the average local price for each year as there are two measurements in one year. Then we need to use the window function and partition data by 'name' of the country and use function `func.lag()` to get the previous value of the 'avg_local_price' column. After that as in the previous solution we find difference between each year and previous one and sort by it.

```
#To perform an operation on a group first,
#we need to partition the data using Window.partitionBy() ,
#and for row number and rank function we need to additionally order by on partition data using Window.orderBy()

window = Window.partitionBy('name').orderBy('name')
```

```
data = datasetDF.selectExpr('name', 'local_price', 'year(date) as year').groupBy('name', 'year')
data = data.withColumn('prev_year_local_price', func.lag('avg_local_price').over(window))
data = data.selectExpr('name', 'year', 'avg_local_price', 'avg_local_price - prev_year_local_price as difference_local_price_from_previous_year')
data.show()
```

name	year	avg_local_price	difference_local_price_from_previous_year
Indonesia	2013	27939.0	4572.0
Indonesia	2018	33625.0	2062.0
Indonesia	2016	30750.0	1530.5
Colombia	2017	9900.0	1500.0
Colombia	2018	11400.0	1500.0
Indonesia	2015	29219.5	1280.5
Colombia	2021	12950.0	1050.0
Indonesia	2020	33500.0	1000.0
Indonesia	2012	23367.0	833.0
Indonesia	2017	31563.0	813.0
Indonesia	2021	34000.0	500.0
Colombia	2016	8400.0	500.0
Colombia	2019	11900.0	500.0
South Korea	2015	4200.0	300.0
Chile	2017	2500.0	300.0
Chile	2021	2965.0	275.0
Chile	2012	2050.0	200.0
South Korea	2016	4350.0	150.0
Argentina	2021	350.0	139.5
Chile	2018	2620.0	120.0

only showing top 20 rows

CPU times: user 29.5 ms, sys: 4.7 ms, total: 34.2 ms

Wall time: 1.15 s

Question 5:

For each country show largest increase in dollar exchange rate (dollar_ex) comparing month of each year and sort in descending order

Saved successfully!



To solve this question, we again use window() and func.lag() to partition data by 'year' and 'name' and to get the previous value of dollar exchange for each row. This time we need to partition by name and year as we need to find previous values for each month separately. For records start from 2011 and for measurements that were made in the first month of each year we get 'NULL' value for func.lag() function as with partitions year and name they don't have dollar_ex for previous rows. Thus, we need to filter resulted data and get rid of rows with null values for the 'previous_dollar_ex' column. Next, we need to find the difference between column 'dollar_ex' with the 'previous_dollar_ex' column. Finally, need to find maximum of their difference grouping by name.

```

window = Window.partitionBy('year', 'name').orderBy('year')
data = datasetDF.selectExpr('name', 'dollar_ex', 'year(date) as year', 'month(date) as month')
data = data.filter('month != 1 and year != 2011').selectExpr('name', 'year', 'dollar_ex', 'ir')
data = data.groupBy('name').agg(F.max('difference_dollar_ex')).sort(desc('max(difference_doll
data.show()

```

name	max(difference_dollar_ex)
Indonesia	1001.0
Colombia	381.6100000000001
South Korea	61.950000000000045
Chile	45.79500000000007
Hungary	24.288050000000027
Pakistan	18.569999999999993
Argentina	11.175049999999999
Sri Lanka	10.5
Russia	9.152499999999996
Japan	9.040000000000006
India	6.584999999999994
Costa Rica	3.9700000000000273
Mexico	3.615500000000001
Philippines	2.839999999999963
Thailand	2.6800000000000033
South Africa	2.281499999999994
Uruguay	1.5050000000000026
Czech Republic	1.3177500000000002
Brazil	1.19855
Turkey	1.086749999999994

only showing top 20 rows

CPU times: user 25.5 ms, sys: 5.17 ms, total: 30.7 ms

Wall time: 1.23 s

Question 6:

Saved successfully!



average dollar_price for BigMac.

This question can be solved by rank() function which helps to identify rank of the value in the partition. So, as we have two measurements for each year, we need to find average for each year. Then we use window function and partition data by name and order by 'avg_dollar_price' column for each year and apply rank() function. Finally, we need to order data by name and rank

```

window = Window.partitionBy('name').orderBy('avg_dollar_price')
data = datasetDF.selectExpr('name', 'dollar_price', 'year(date) as year').groupBy('name', 'year')
data = data.withColumn('rank', rank().over(window)).sort('name', 'rank')

```



```
data.show()
```

```
+-----+-----+-----+-----+
|   name|year| avg_dollar_price|rank|
+-----+-----+-----+-----+
|Argentina|2019|2.4364536274995503| 1|
|Argentina|2014|2.8024533848179196| 2|
|Argentina|2016|2.8687715986163003| 3|
|Argentina|2015|3.1585833099710303| 4|
|Argentina|2020| 3.178059921365675| 5|
|Argentina|2018| 3.332767902579255| 6|
|Argentina|2017| 3.796962230447655| 7|
|Argentina|2021| 3.846425460025445| 8|
|Argentina|2013| 3.848925979818775| 9|
|Argentina|2012| 4.398784797986675|10|
|Argentina|2011| 4.83968542044767|11|
|Australia|2016| 4.024196250690875| 1|
|Australia|2015| 4.120484999455475| 2|
|Australia|2019| 4.306075000000005| 3|
|Australia|2017| 4.40156750080276| 4|
|Australia|2020| 4.514797500000001| 5|
|Australia|2018| 4.610776249831865| 6|
|Australia|2014| 4.64384374945093| 7|
|Australia|2013| 4.76015962477957| 8|
|Australia|2012| 4.808357999562684| 9|
+-----+-----+-----+-----+
```

only showing top 20 rows

▼ Execution Time:

Question	cluster,sec	local,sec
1	7.91	7.21
2	11.31	10.96
3	11.25	10.9
4	17.36	21.6
	14.36	16.36
	12.63	16.23
all together	26.677	43.617

Saved successfully!



Enter to this directory path:

user_lsc_25/project/project_1.py

For runing the file locally use:

spark-submit --master local project_1.py

For running the file through yarn:

```
spark-submit --master yarn project_1.py
```

Insight:

Running the project file in the local mode was faster for question 1, 2 and 3, which has simple queries with some groupby, filter and aggregation functions. We assume that clustering is slower because it spends some time for scheduling and assigning task to each machine. Local mode right away working on the code. Starting from the 4th Question, we have started using partitioning for solutions and we get more complex queries. As a result, we get more faster results in yarn mode as it gives opportunity to run processes in parallel. Running all question implementations is almost two times faster in cluster mode because it distributes all the task to clusters and run it in parallel. Resultingly, it's better to use cluster mode for complex queries and local for simple ones.

Double-click (or enter) to edit

References:

[1]"The Big Mac Economic Index", Kaggle.com, 2022. [Online]. Available:

<https://www.kaggle.com/yamqwe/the-big-mac-economic-index>. [Accessed: 18- Jan- 2022].

[2]"Spark SQL and DataFrames - Spark 3.2.0 Documentation", Spark.apache.org, 2022. [Online].

Available: <https://spark.apache.org/docs/latest/sql-programming-guide.html>. [Accessed: 18- Jan- 2022].

[3]"Apache Spark DS DF and RDD",

<https://docs.google.com/presentation/d/194AYzBioTdgcgczpZmBEjmQ6fe6OT0z/edit#slide=id.p3>, 2020.

Saved successfully!





Saved successfully! 